


Transformers

CoE197Z/EE298Z (Deep Learning)

Rowel Atienza, Ph.D.

rowel@eee.upd.edu.ph

Encoder-Decoder Sequence-to-Sequence



seq2seq

RNN

Serial

Difficult to parallelize

Uni-directional

 Bi-directional version is
 much slower

Susceptible to catastrophic
forgetting

Slow

Transformer


Parallel

Bidirectional


Not susceptible to catastrophic
forgetting

Fast

Transformer



Language Translation




Transformer Encoder Unit Details

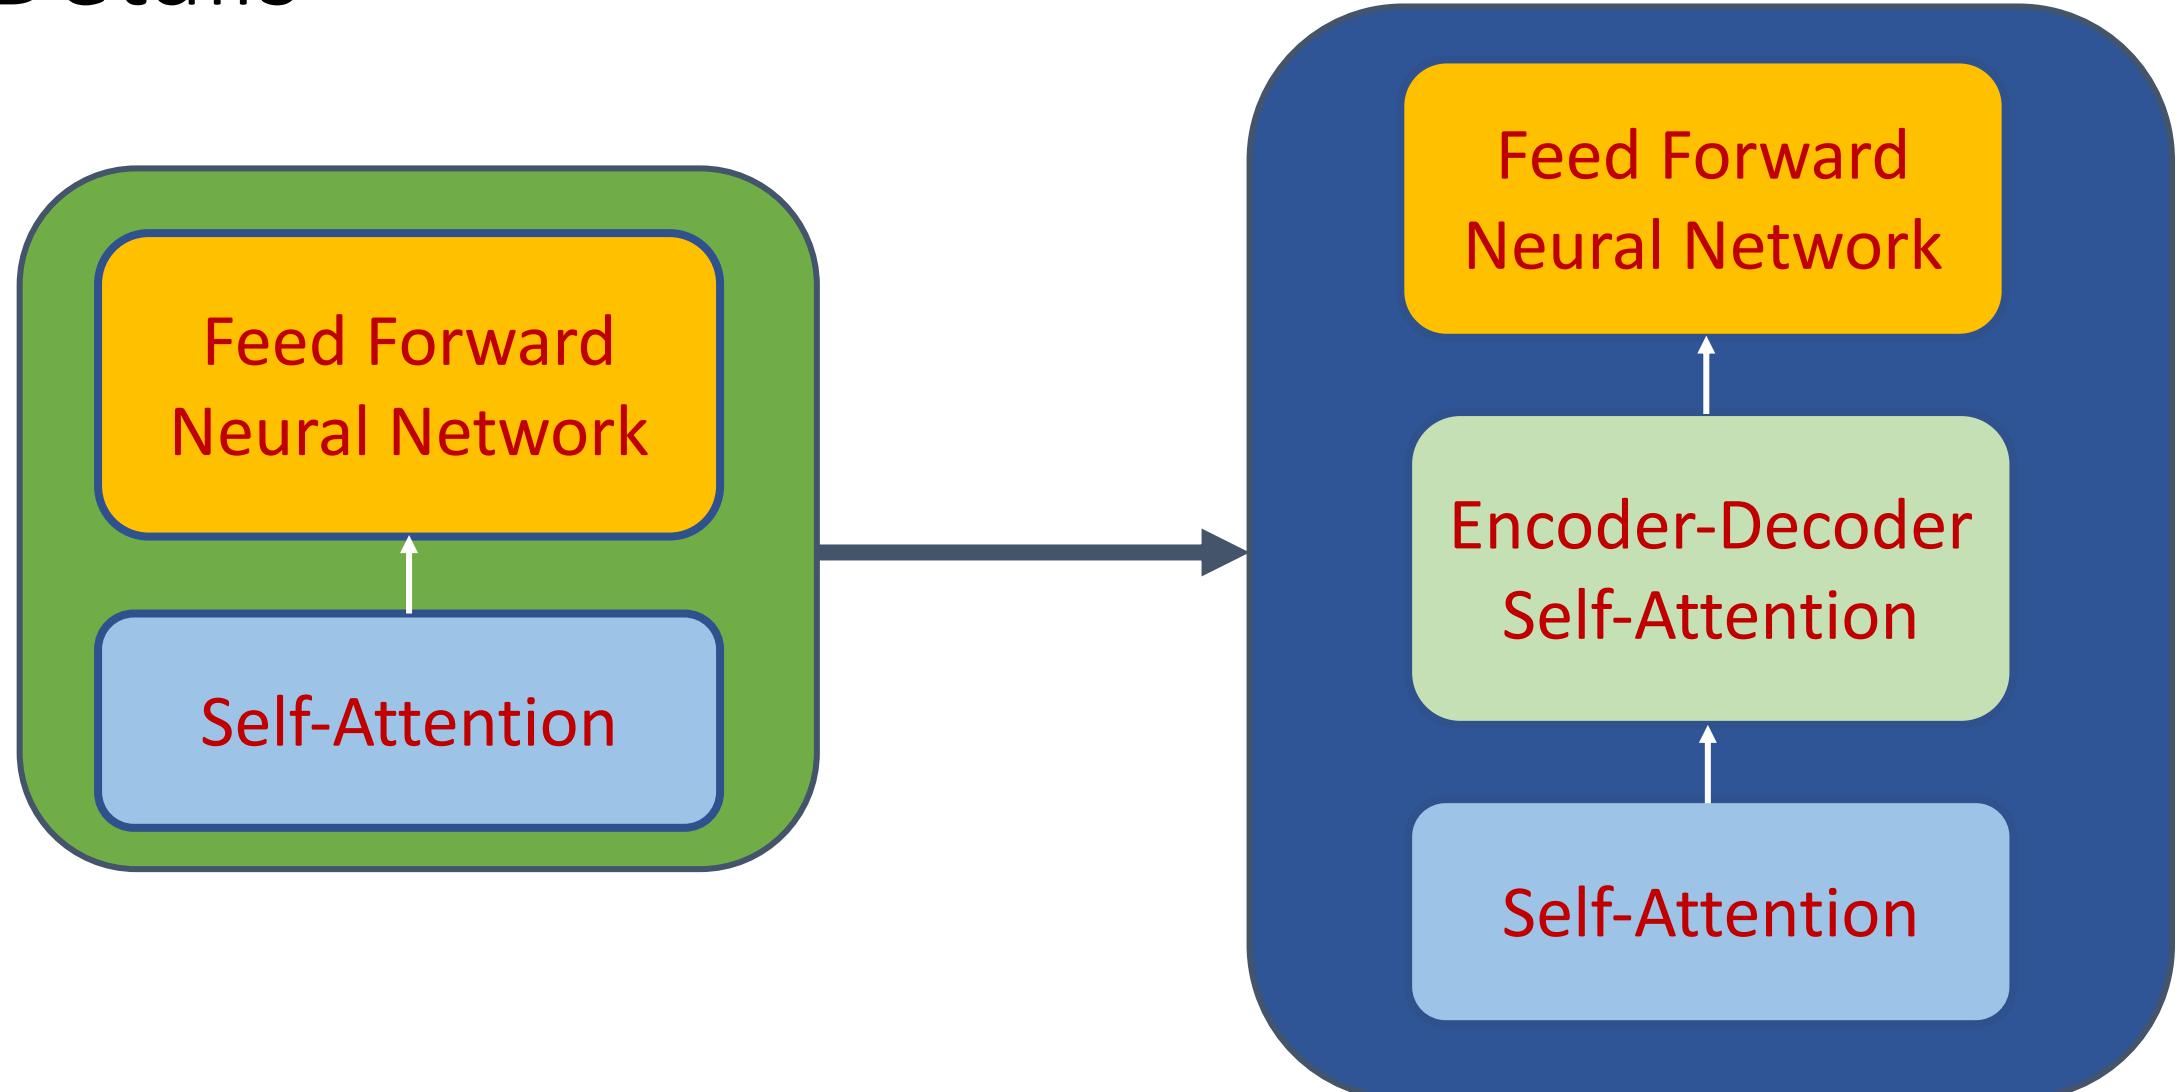
No recurrence (No RNN)

No CNN

Operations: Linear, Norm, Matrix
Multiply, Dot Product, Softmax



Transformer Encoder and Decoder Unit Details



Input Embedding is an $n - \text{dim}$ vector

$$\boldsymbol{x}_1^T$$

.1	-2	.4	-1
----	----	----	----

tayo

$$\boldsymbol{x}_2^T$$

.3	1	-1	2
----	---	----	---

ay


$$\boldsymbol{x}_3^T$$

.1	.0	1	-1
----	----	---	----


masaya

Example: Each word is converted into a 512-dim embedding vector.
In the simple example above, it is 4-dim.

Encoder
(Typically 6
Encoder Units)

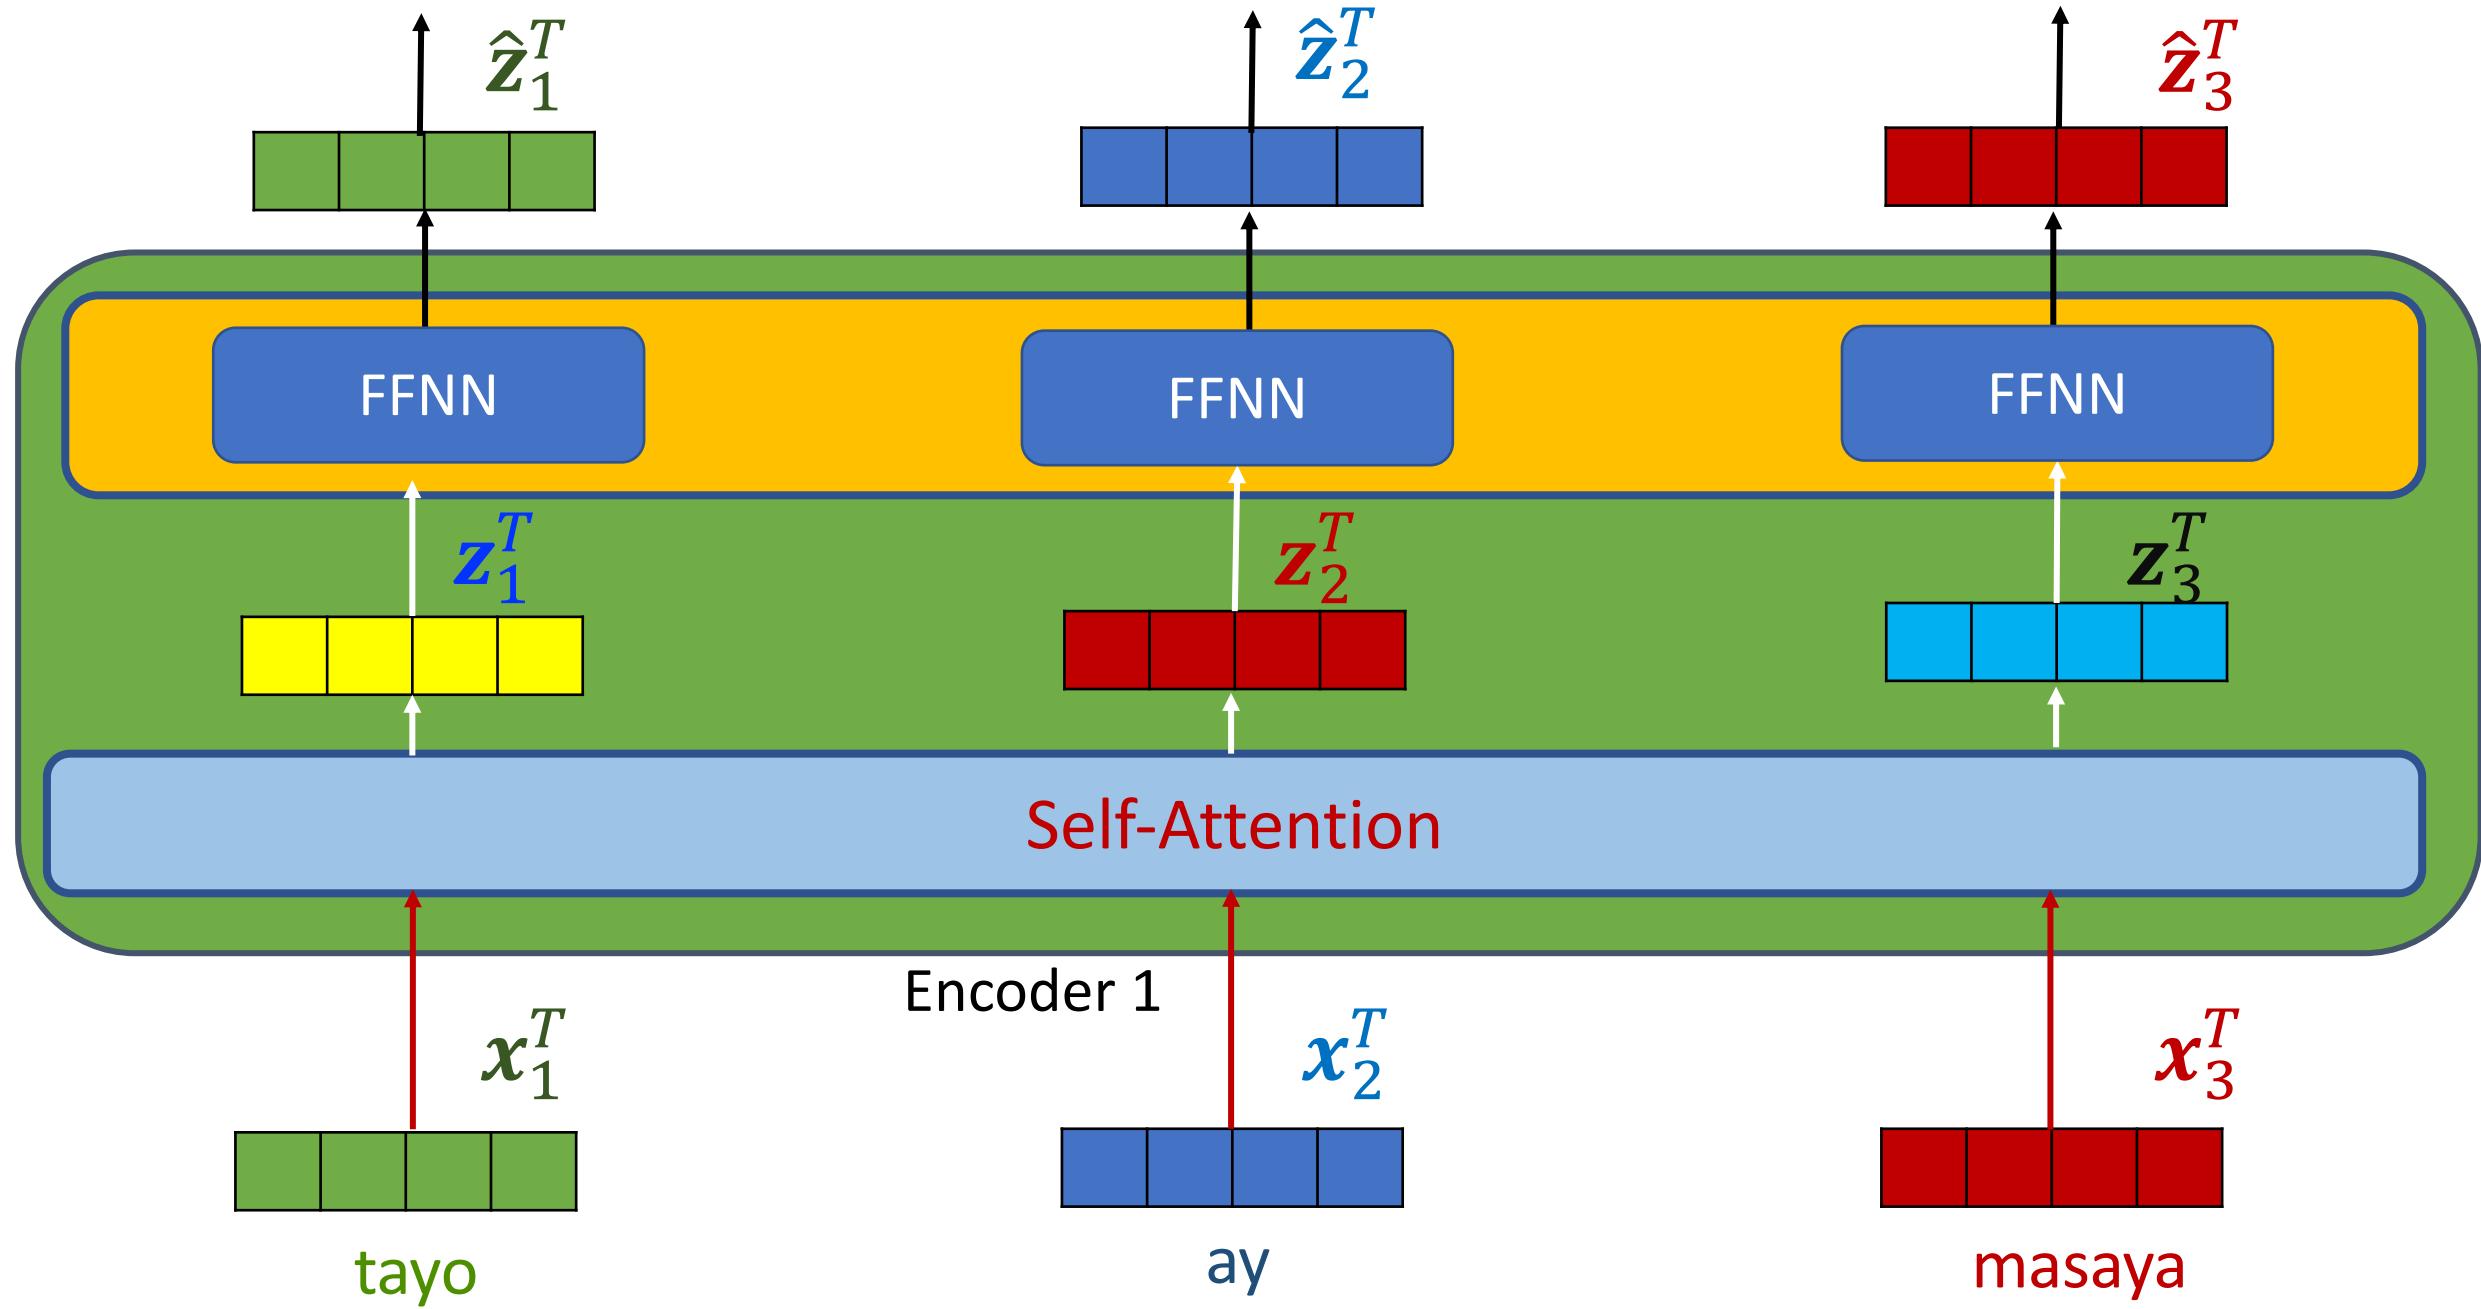


The Length of the Input is n


 x_1^T  $word_1$ x^T_{\dots} \dots x_n^T  $word_n$

Example: n could be the maximum possible length of a sentence.

Encoder with Latent Variables \mathbf{z}_i



Attention between 2 words



Attention as measured by the width of the arrow

Attention Layer 1 Learnable Parameters

Self-Attention

$$W^Q$$



$$\mathbf{q}_1^T = \boldsymbol{x}_1^T W^Q$$

$$\mathbf{q}_2^T = \boldsymbol{x}_2^T W^Q$$

$$\mathbf{q}_3^T = \boldsymbol{x}_3^T W^Q$$

Queries

$$Q = XW^Q$$

$$W^K$$

$$\mathbf{k}_1^T = \boldsymbol{x}_1^T W^K$$

$$\mathbf{k}_2^T = \boldsymbol{x}_2^T W^K$$

$$\mathbf{k}_3^T = \boldsymbol{x}_3^T W^K$$

Keys

$$K = XW^K$$

$$W^V$$

$$\mathbf{v}_1^T = \boldsymbol{x}_1^T W^V$$

$$\mathbf{v}_2^T = \boldsymbol{x}_2^T W^V$$

$$\mathbf{v}_3^T = \boldsymbol{x}_3^T W^V$$

Values


$$V = XW^V$$

$$X = \begin{bmatrix} \boldsymbol{x}_1^T \\ \boldsymbol{x}_2^T \\ \boldsymbol{x}_3^T \end{bmatrix}$$

Encoder 1
Inputs

$$X = \begin{bmatrix} \textcolor{green}{x}_1^T \\ \textcolor{blue}{x}_2^T \\ \textcolor{red}{x}_3^T \end{bmatrix} \quad \begin{array}{l} \text{Encoder 1} \\ \text{Inputs} \end{array}$$


Queries



$$Q$$

$$Q = XW^Q$$

Keys



$$K = XW^K$$

Scores

$$\begin{aligned} & \left(\begin{array}{c|c} \textcolor{purple}{X} & W^Q \\ \hline \textcolor{yellow}{X} & W^K \end{array} \right)^T \\ &= \begin{bmatrix} \textcolor{purple}{X}^T \\ \textcolor{yellow}{X}^T \end{bmatrix} \left(\begin{bmatrix} W^Q \\ W^K \end{bmatrix} \right)^T = S \end{aligned}$$

tayo $\textcolor{green}{x}_1^T$




ay $\textcolor{blue}{x}_2^T$



masaya $\textcolor{red}{x}_3^T$



W^Q




$$\textcolor{purple}{q}_1^T = \textcolor{green}{x}_1^T W^Q$$

$$\textcolor{purple}{q}_2^T = \textcolor{blue}{x}_2^T W^Q$$

$$\textcolor{purple}{q}_3^T = \textcolor{red}{x}_3^T W^Q$$

W^K



$$\textcolor{blue}{k}_1^T = \textcolor{green}{x}_1^T W^K$$

$$\textcolor{blue}{k}_2^T = \textcolor{blue}{x}_2^T W^K$$

$$\textcolor{blue}{k}_3^T = \textcolor{red}{x}_3^T W^K$$

$$\begin{bmatrix} s_{11} = \textcolor{purple}{q}_1^T \textcolor{blue}{k}_1 \\ s_{12} = \textcolor{purple}{q}_1^T \textcolor{blue}{k}_2 \\ s_{13} = \textcolor{purple}{q}_1^T \textcolor{blue}{k}_3 \end{bmatrix}^T$$

$$\begin{bmatrix} s_{21} = \textcolor{purple}{q}_2^T \textcolor{blue}{k}_1 \\ s_{22} = \textcolor{purple}{q}_2^T \textcolor{blue}{k}_2 \\ s_{23} = \textcolor{purple}{q}_2^T \textcolor{blue}{k}_3 \end{bmatrix}^T$$

$$\begin{bmatrix} s_{31} = \textcolor{purple}{q}_3^T \textcolor{blue}{k}_1 \\ s_{32} = \textcolor{purple}{q}_3^T \textcolor{blue}{k}_2 \\ s_{33} = \textcolor{purple}{q}_3^T \textcolor{blue}{k}_3 \end{bmatrix}^T$$

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

d_k is keys/queries dim (e.g. 4)


$$\text{Attention} = \text{softmax} \left(\frac{\begin{array}{c} \text{purple matrix} \\ \text{yellow matrix}^T \end{array}}{\sqrt{d_k}} \right) \text{yellow matrix}$$
$$\text{Attention}(Q, K, V) = Z = \text{green matrix}$$


Values When all things considered, what my outputs should be

Keys What others think my outputs should be

Queries What I think my outputs should be


Attention = *softmax*

Attention(*Q*, *K*, *V*) = *Z* = 




Consider an Attention Layer Examining a Digit

Quiet! I can see everything that you can see. My friend, you are a part of digit 3.




Example: Let us focus on the lower-right patch only

Multi-Head
(eg 8-head)




Multi-Head
(eg 8-head)




Multi-Head
(eg 8-head)
Merge Outputs
Apply Weights

Encoder 1
Inputs



$$X = \begin{bmatrix} x_1^T \\ x_2^T \\ x_3^T \end{bmatrix}$$


Self-Attention




$$cat(Z_1, \dots, Z_8)$$

$$= Z \quad \begin{smallmatrix} \text{---} \\ \text{---} \\ \text{---} \end{smallmatrix}$$

$$\times \quad W^O$$



Adding Position Info to Inputs



Positional Encoding

$$PE_{(pos,2i)} = \sin\left(\frac{pos}{10000^{\frac{2i}{d_k}}}\right) \quad dim = 2i \text{ is even}$$

$$PE_{(pos,2i+1)} = \cos\left(\frac{pos}{10000^{\frac{2i+1}{d_k}}}\right) \quad dim = 2i + 1 \text{ is odd}$$

$pos = 0, 1, \dots, n_{pos-1}$

$dim = 0, 1, \dots, n_{dim-1}$

Other positional encoding methods: learnable

Assuming $n_{pos-1} = 2, n_{dim-1} = 3, d_k = 4$


<i>pos</i>	<i>dim</i>			
	0	1	2	3
0	$\sin\left(\frac{0}{10000^{0/4}}\right)$	$\cos\left(\frac{0}{10000^{0/4}}\right)$	$\sin\left(\frac{0}{10000^{2/4}}\right)$	$\cos\left(\frac{0}{10000^{2/4}}\right)$
1	$\sin\left(\frac{1}{10000^{0/4}}\right)$	$\cos\left(\frac{1}{10000^{0/4}}\right)$	$\sin\left(\frac{1}{10000^{2/4}}\right)$	$\cos\left(\frac{1}{10000^{2/4}}\right)$
2	$\sin\left(\frac{2}{10000^{0/4}}\right)$	$\cos\left(\frac{2}{10000^{0/4}}\right)$	$\sin\left(\frac{2}{10000^{2/4}}\right)$	$\cos\left(\frac{2}{10000^{2/4}}\right)$

Improvements:

Residual
Connections


Layer Norm


Encoder 1




FFN: Feed Forward Neural Network (MLP)

$$FFN(x) = \max(0, xW_1 + b_1)W_2 + b_2$$







*Masking prevents Decoder 1 from seeing the future.
Decoder 1 relies only on the previous outputs.*

Vision Transformer




Figure 1: Model overview. We split an image into fixed-size patches, linearly embed each of them, add position embeddings to the resulting sequence of vectors, and feed the patches to a standard Transformer encoder. In order to perform classification, we use the standard approach of adding an extra learnable "classification token" to the sequence. The illustration of the Transformer encoder was inspired by [Vaswani et al. \(2017\)](#).

AN IMAGE IS WORTH 16X16 WORDS:
 TRANSFORMERS FOR IMAGE RECOGNITION AT SCALE, ICLR 2021 Submission


```
class ViT(nn.Module):
    def __init__(self, *, image_size, patch_size, num_classes, dim, depth, heads, mlp_dim, channels = 3):
        super().__init__()
        assert image_size % patch_size == 0, 'image dimensions must be divisible by the patch size'
        num_patches = (image_size // patch_size) ** 2
        patch_dim = channels * patch_size ** 2

        self.patch_size = patch_size

        self.pos_embedding = nn.Parameter(torch.randn(1, num_patches + 1, dim))
        self.patch_to_embedding = nn.Linear(patch_dim, dim)
        self.cls_token = nn.Parameter(torch.randn(1, 1, dim))
        self.transformer = Transformer(dim, depth, heads, mlp_dim)


        self.to_cls_token = nn.Identity()
```

$28 \times 28 \times 1$



$(P, P) = (14, 14)$


4×196




3.1 VISION TRANSFORMER (ViT)

Our Transformer for images follows the architecture designed for NLP. Figure 1 depicts the setup. The standard Transformer receives as input a 1D sequence of token embeddings. To handle 2D images, we reshape the image $\mathbf{x} \in \mathbb{R}^{H \times W \times C}$ into a sequence of flattened 2D patches $\mathbf{x}_p \in \mathbb{R}^{N \times (P^2 \cdot C)}$. (H, W) is the resolution of the original image and (P, P) is the resolution of each image patch. $N = HW/P^2$ is then the effective sequence length for the Transformer. The Transformer uses constant widths through all of its layers, so a trainable linear projection maps each vectorized patch to the model dimension D (Eq. 1), the output of which we refer to as our patch embeddings.

$28 \times 28 \times 1$




4×196



$(P, P) = (14, 14)$

```
def forward(self, img, mask = None):
    p = self.patch_size

    x = rearrange(img, 'b c (h p1) (w p2) -> b (h w) (p1 p2 c)', p1 = p, p2 = p)
```



```


self.patch_to_embedding = nn.Linear(patch_dim, dim)

def forward(self, img, mask = None):
    p = self.patch_size


    x = rearrange(img, 'b c (h p1) (w p2) -> b (h w) (p1 p2 c)', p1 = p, p2 = p)
    x = self.patch_to_embedding(x)

```

Patch
 4×196



Embedding
 4×128



`self.patch_to_embedding = nn.Linear(patch_dim, dim)`

A red bracket is shown under the first three patch features (blue, orange, and red). A red arrow points from this bracket upwards towards the embedding stage, indicating that these features are processed by a linear layer before being combined with the fourth patch feature and the position embedding.

$$\mathbf{z}_0 = [\mathbf{x}_{\text{class}}; \mathbf{x}_p^1 \mathbf{E}; \mathbf{x}_p^2 \mathbf{E}; \dots; \mathbf{x}_p^N \mathbf{E}] + \mathbf{E}_{pos}, \quad \mathbf{E} \in \mathbb{R}^{(P^2 \cdot C) \times D}, \mathbf{E}_{pos} \in \mathbb{R}^{(N+1) \times D}$$

Class Token

1×128



Similar to BERT's [class] token, we prepend a learnable embedding to the sequence of embedded patches ($\mathbf{z}_0^0 = \mathbf{x}_{\text{class}}$), whose state at the output of the Transformer encoder (\mathbf{z}_0^L) serves as the

```
self.cls_token = nn.Parameter(torch.randn(1, 1, dim))

def forward(self, img, mask = None):
    p = self.patch_size

    x = rearrange(img, 'b c (h p1) (w p2) -> b (h w) (p1 p2 c)', p1 = p, p2 = p)
    x = self.patch_to_embedding(x)

    cls_tokens = self.cls_token.expand(img.shape[0], -1, -1)
```

x Embedding 4×128



Class Token 1×128



x 5×128



`x = torch.cat((cls_tokens, x), dim=1)`

```
def forward(self, img, mask = None):
    p = self.patch_size
    x = rearrange(img, 'b c (h p1) (w p2) -> b (h w) (p1 p2 c)', p1 = p, p2 = p)
    x = self.patch_to_embedding(x)


    cls_tokens = self.cls_token.expand(img.shape[0], -1, -1)
    x = torch.cat((cls_tokens, x), dim=1)
```

$$\mathbf{z}_0 = [\mathbf{x}_{\text{class}}; \mathbf{x}_p^1 \mathbf{E}; \mathbf{x}_p^2 \mathbf{E}; \dots; \mathbf{x}_p^N \mathbf{E}] + \mathbf{E}_{\text{pos}},$$

Position Embedding


Position embeddings are added to the patch embeddings to retain positional information. We explore different 2D-aware variants of position embeddings (Appendix C.3) without any significant gains over standard 1D position embeddings. The joint embedding serves as input to the encoder.


```
self.pos_embedding = nn.Parameter(torch.randn(1, num_patches + 1, dim))
```

x 5×128  $+$ $\text{Position Embedding}$ 5×128  x 5×128 $=$ 


$$\mathbf{z}_0 = [\mathbf{x}_{\text{class}}; \mathbf{x}_p^1 \mathbf{E}; \mathbf{x}_p^2 \mathbf{E}; \dots; \mathbf{x}_p^N \mathbf{E}] + \mathbf{E}_{pos}, \quad \mathbf{E} \in \mathbb{R}^{(P^2 \cdot C) \times D}, \mathbf{E}_{pos} \in \mathbb{R}^{(N+1) \times D}$$

$x \quad \textcolor{blue}{+=} \quad \text{self.pos_embedding}$





```
x = self.to_cls_token(x[:, 0])
return self.mlp_head(x)
```



```
self.mlp_head = nn.Sequential(
    nn.Linear(dim, mlp_dim),
    nn.GELU(),
    nn.Linear(mlp_dim, num_classes)
)
```

image representation y (Eq. 4). Both during pre-training and fine-tuning, the classification head is attached to \mathbf{z}_L^0 .

```
x = self.to_cls_token(x[:, 0])
return self.mlp_head(x)
```

```
def forward(self, img, mask = None):
    p = self.patch_size

    x = rearrange(img, 'b c (h p1) (w p2) -> b (h w) (p1 p2 c)', p1 = p, p2 = p)
    x = self.patch_to_embedding(x)

    cls_tokens = self.cls_token.expand(img.shape[0], -1, -1)
    x = torch.cat((cls_tokens, x), dim=1)
    x += self.pos_embedding
    x = self.transformer(x, mask)

    x = self.to_cls_token(x[:, 0])
    return self.mlp_head(x)
```

Transformer

```
class Transformer(nn.Module):
    def __init__(self, dim, depth, heads, mlp_dim):
        super().__init__()
        self.layers = nn.ModuleList([])
        for _ in range(depth):
            self.layers.append(nn.ModuleList([
                Residual(PreNorm(dim, Attention(dim, heads = heads))),
                Residual(PreNorm(dim, FeedForward(dim, mlp_dim)))
            ]))
    def forward(self, x, mask = None):
        for attn, ff in self.layers:
            x = attn(x, mask = mask)
            x = ff(x)
        return x
```

$$\begin{aligned} \mathbf{z}'_\ell &= \text{MSA}(\text{LN}(\mathbf{z}_{\ell-1})) + \mathbf{z}_{\ell-1}, & \ell &= 1 \dots L \\ \mathbf{z}_\ell &= \text{MLP}(\text{LN}(\mathbf{z}'_\ell)) + \mathbf{z}'_\ell, & \ell &= 1 \dots L \\ \mathbf{y} &= \text{LN}(\mathbf{z}_L^0) \end{aligned}$$

Residual

```
class Residual(nn.Module):
    def __init__(self, fn):
        super().__init__()
        self.fn = fn
    def forward(self, x, **kwargs):
        return self.fn(x, **kwargs) + x
```

Layer Norm

```
class PreNorm(nn.Module):
    def __init__(self, dim, fn):
        super().__init__()
        self.norm = nn.LayerNorm(dim)
        self.fn = fn
    def forward(self, x, **kwargs):
        return self.fn(self.norm(x), **kwargs)
```

Feed Forward (MLP)

```
class FeedForward(nn.Module):
    def __init__(self, dim, hidden_dim):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(dim, hidden_dim),
            nn.GELU(),
            nn.Linear(hidden_dim, dim)
        )
    def forward(self, x):
        return self.net(x)
```

Attention

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

```
class Attention(nn.Module):
    def __init__(self, dim, heads = 8):
        super().__init__()
        self.heads = heads
        self.scale = dim ** -0.5


        self.to_qkv = nn.Linear(dim, dim * 3, bias = False)
        self.to_out = nn.Linear(dim, dim)

    def forward(self, x, mask = None):
        b, n, _, h = *x.shape, self.heads
        qkv = self.to_qkv(x)
        q, k, v = rearrange(qkv, 'b n (qkv h d) -> qkv b h n d', qkv = 3, h = h)
        dots = torch.einsum('bhid,bhjd->bhij', q, k) * self.scale
```

$$\frac{QK^T}{\sqrt{d_k}}$$


Multi-Head
(eg 8-head)
Merge Outputs
Apply Weights

Encoder 1
Inputs



$$X = \begin{bmatrix} x_1^T \\ x_2^T \\ x_3^T \end{bmatrix}$$


Self-Attention



$$cat(Z_1, \dots, Z_8)$$

$$= Z \quad \begin{smallmatrix} \text{---} \\ \text{---} \\ \text{---} \end{smallmatrix}$$

$$\times \quad W^O$$



$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

```
attn = dots.softmax(dim=-1)
```

```
out = torch.einsum('bhij,bhjd->bhid', attn, v)
```

```
out = rearrange(out, 'b h n d -> b n (h d)')
```

```
out = self.to_out(out)
```

```
return out
```

$\times W^O = z$

$cat(z_1, \dots, z_8)$

Function composition

$$f_n(\cdot)$$

MLP Head

$$f_6(\cdot)$$


$$f_1(x)$$

Encoder 6

Encoder 1

...

$$28 \times 28 \times 1$$



$$(P, P) = (14, 14)$$

Inductive Bias

Transformers lack some inductive biases inherent to CNNs, such as translation equivariance and locality, and therefore do not generalize well when trained on insufficient amounts of data.

However, the picture changes if we train the models on large datasets (14M-300M images). We find that large scale training trumps inductive bias.

CNN Model on MNIST: ~99.2% 15mins to train on CPU

Transformer Model on MNIST: ~98.2% 7mins to train on CPU

References

Vaswani, Ashish, et al. "Attention is all you need." *Advances in neural information processing systems*. 2017.

Illustrated Transformer, <http://jalammar.github.io/illustrated-transformer/>

Transformers from Scratch, <http://peterbloem.nl/blog/transformers>

Transformer Family, <https://lilianweng.github.io/lil-log/2020/04/07/the-transformer-family.html>

<https://github.com/lucidrains/vit-pytorch>

In Summary

Transformer could be the most important breakthrough in the recent history of deep learning

Transformer has been used to produce state-of-the-art performances in NLP and vision

Expect more development in this field in the near future




Figure 1: The Transformer - model architecture.

Vaswani, Ashish, et al. "Attention is all you need." *Advances in neural information processing systems*. 2017.